

Jon Saad-Falcon

jonsaadfalcon@stanford.edu | jonsaadfalcon.com | [Google Scholar](#)

Scaling Intelligence Lab and Hazy Research, Stanford University

Education

- 2023–present **Ph.D.**, Computer Science, Stanford University
Advisors: Christopher Ré and Azalia Mirhoseini
- 2023–present **M.B.A.**, Stanford Graduate School of Business
- 2018–2022 **B.S.–M.S.**, Computer Science, Georgia Institute of Technology
Advisor: Duen Horng (Polo) Chau. *Minors:* Mathematics and Linguistics

Research Experience

- 2023–present **PhD Researcher**, Stanford University, Computer Science
Areas: language models, ML systems, intelligence efficiency, inference-time techniques.
- 2026–present **Student Researcher**, Google — TPU and Gemini teams (Mountain View, CA, USA)
- 2023 **Research Intern**, Databricks, ML Ops Team (San Francisco, CA, USA)
With Prof. Matei Zaharia. Built an automated evaluation system for retrieval-augmented generation.
- 2022–2023 **Fulbright Scholar**, Humboldt-Universität zu Berlin (Berlin, Germany)
With Prof. Ulf Leser. Developed AI-powered techniques for patient treatment design using electronic health records.
- 2021–2022 **Predoctoral Young Investigator**, Allen Institute for AI (AI2) — Semantic Scholar (Seattle, WA, USA)
Mentors: Prof. Doug Downey and Prof. Daniel Weld. Efficiency techniques for caching sequence representations in language models.
- 2021 **Research Intern**, Stanford University, CSLI (Stanford, CA, USA)
With Prof. Christopher Potts and Prof. Matei Zaharia. NSF-funded REU; built the LoTTE benchmark and contributed to ColBERTv2.
- 2019–2022 **Undergraduate Research Assistant**, Georgia Institute of Technology (Atlanta, GA, USA)
Advisors: Prof. Duen Horng (Polo) Chau and Prof. Diyi Yang.
Projects in NLP and human-computer interaction.

Awards & Honors

- 2025 **JP Morgan AI/ML Fellowship**
- 2023 **Stanford Graduate Fellowship**
- 2023 **Stanford EDGE Fellowship**
- 2023 **GEM PhD Fellowship**, *National GEM Consortium*
- 2023 **Knight-Hennessy Scholarship**, Finalist, *Stanford University*
- 2023 **Gates-Cambridge Scholarship**, *Bill & Melinda Gates Foundation*
- 2023 **College of Engineering Fellowship**, *University of California, Berkeley*
- 2022 **Fulbright Scholarship**, Research Award, *U.S. State Department*
- 2022 **Summer Venture in Management Program (SVMP)**, *Harvard Business School*

2021	Donald V. Jackson Fellowship , <i>Georgia Tech</i>
2021	Computer Science Research Mentorship Program , <i>Google</i>
2021	U.N. Millennium Fellowship , <i>United Nations</i>
2020	D.E. Shaw Nexus Fellowship
2020	President’s Undergraduate Research Award , <i>Georgia Tech</i>
2018	Stamps President’s Scholarship , <i>Georgia Tech & Stamps Foundation</i>

Publications

Preprints

- P1. Saad-Falcon, J.**, Narayan, A., Akengin, O., Shandilya, H., Manihani, R., Bo, G., Hennessy, J., Mirhoseini, A. & Ré, C. *OpenJarvis: Personal AI, On Personal Devices*. 2026. [Blog](#) | [Code](#).
- P2. Saad-Falcon, J.**, Narayan, A., Akengin, O., Griffin, W., Shandilya, H., Lafuente, A., Goel, M., Joseph, R., Natarajan, S., Guha, E., Zhu, S., Athiwaratkun, B., Hennessy, J., Mirhoseini, A. & Ré, C. *Intelligence per Watt: Measuring Intelligence Efficiency of Local AI*. 2025. arxiv.org/abs/2511.07885.
- P3.** Kang, H., Suresh, T., **Saad-Falcon, J.** & Mirhoseini, A. *TRACE: Capability-Targeted Agentic Training*. 2026. arxiv.org/abs/2604.05336.

Peer-reviewed Conference Proceedings

- C1. Saad-Falcon, J.**, Buchanan, E. K., Chen, M. F., Huang, T.-H., McLaughlin, B., Bhathal, T., Zhu, S., Athiwaratkun, B., Sala, F., Linderman, S., Mirhoseini, A. & Ré, C. *Weaver: Shrinking the Generation-Verification Gap with Weak Verifiers* in *Advances in Neural Information Processing Systems (NeurIPS)* (2025). arxiv.org/abs/2506.18203.
- C2. Saad-Falcon, J.**, Vivek, R., Berrios, W., Naik, N., Franklin, M., Vidgen, B., Singh, A., Kiela, D. & Mehri, S. *LMUnit: Fine-grained Evaluation with Natural Language Unit Tests* in *Proceedings of the 2025 Conference on Empirical Methods in Natural Language Processing (EMNLP)* (2025). arxiv.org/abs/2412.13091.
- C3. Saad-Falcon, J.**, Lafuente Gamarra, A., Natarajan, S., Maru, N., Todorov, H., Guha, E., Buchanan, E. K., Chen, M., Guha, N., Ré, C. & Mirhoseini, A. *Archon: An Architecture Search Framework for Inference-Time Techniques* in *International Conference on Machine Learning (ICML) ★ Oral*, ICLR 2025 SSI-FM Workshop (2025). arxiv.org/abs/2409.15254.
- C4. Saad-Falcon, J.**, Fu, D. Y., Arora, S., Guha, N. & Ré, C. *Benchmarking and Building Long-Context Retrieval Models with LoCo and M2-BERT* in *International Conference on Machine Learning (ICML)* (2024). arxiv.org/abs/2402.07440.
- C5. Saad-Falcon, J.**, Barrow, J., Siu, A., Nenkova, A., Yoon, S., Rossi, R. A. & Dernoncourt, F. *PDFTriage: Question Answering over Long, Structured Documents* in *Proceedings of EMNLP 2024: Industry Track* (2024). aclanthology.org/2024.emnlp-industry.13.
- C6. Saad-Falcon, J.**, Khattab, O., Potts, C. & Zaharia, M. *ARES: An Automated Evaluation Framework for Retrieval-Augmented Generation Systems* in *Proceedings of NAACL ★ Oral* (2024). arxiv.org/abs/2311.09476.
- C7. Saad-Falcon, J.**, Khattab, O., Santhanam, K., Florian, R., Roukos, S., Sultan, M. A., Sil, A., Zaharia, M. & Potts, C. *UDAPDR: Unsupervised Domain Adaptation via LLM Prompting and Distillation of Rerankers* in *Proceedings of EMNLP* (2023). arxiv.org/abs/2303.00807.
- C8. Saad-Falcon, J.**, Singh, A., Soldaini, L., D’Arcy, M., Cohan, A. & Downey, D. *Embedding Recycling for Language Models* in *Findings of EACL* (2023). arxiv.org/abs/2207.04993.
- C9.** Santhanam*, K., **Saad-Falcon*, J.**, Franz, M., Khattab, O., Sil, A., Florian, R., Sultan, M. A., Roukos, S., Zaharia, M. & Potts, C. *Moving Beyond Downstream Task Accuracy for Information Retrieval Benchmarking* in *Findings of ACL* (2023).

- C10.** Lahav, D., **Saad-Falcon, J.**, Kuehl, B., Johnson, S., Parasa, S., Shomron, N., Chau, D. H., Yang, D., Horvitz, E., Weld, D. S. & Hope, T. *A Search Engine for Discovery of Scientific Challenges and Directions* in AAAI ★ **Oral** (2022). arxiv.org/abs/2108.13751.
- C11.** Santhanam*, K., Khattab*, O., **Saad-Falcon, J.**, Potts, C. & Zaharia, M. *ColBERTv2: Effective and Efficient Retrieval via Lightweight Late Interaction* in *Proceedings of NAACL* (2022). arxiv.org/abs/2112.01488.
- C12.** Li, K., Yang, H., Upadhayay, A., Zhou, Z., **Saad-Falcon, J.** & Chau, D. H. *Argo Scholar: Interactive Visual Exploration of Literature in Browsers* in *IEEE VIS* ★ **Best Poster, Honorable Mention** (2021).
- C13.** Shaikh, O., **Saad-Falcon, J.**, Wright, A. P., Das, N., Freitas, S., Asensio, O. I. & Chau, D. H. *EnergyVis: Interactively Tracking and Exploring Energy Consumption for ML Models* in *CHI Late-Breaking Work* (2021). arxiv.org/abs/2103.16435.
- C14.** Shaikh, O., Chen, J., **Saad-Falcon, J.**, Chau, D. H. & Yang, D. *Examining the Ordering of Rhetorical Strategies in Persuasive Requests* in *Findings of EMNLP* (2020). arxiv.org/abs/2010.04625.
- C15.** **Saad-Falcon, J.**, Shaikh, O., Wang, Z. J., Wright, A. P., Richardson, S. & Chau, D. H. *Mapping Researchers with PeopleMap* in *IEEE VIS* ★ **Best Poster, Honorable Mention** (2020). arxiv.org/abs/2009.00091.

Invited Talks

- T1.** **Saad-Falcon, J.** *Intelligence per Watt: Measuring Intelligence Efficiency of Local AI*. EE292P: Atoms, Bits, and National Interest, Stanford University. Feb. 2026.
- T2.** **Saad-Falcon, J.** *Intelligence per Watt: Measuring Intelligence Efficiency of Local AI*. Two Sigma. Jan. 2026.
- T3.** **Saad-Falcon, J.** *Intelligence per Watt: Measuring Intelligence Efficiency of Local AI*. MBZUAI Speech and NLP Symposium. Jan. 2026.
- T4.** **Saad-Falcon, J.** *Weaver: Shrinking the Generation-Verification Gap with Weak Verifiers*. Measuring Intelligence Summit, PyTorch Conference. Oct. 2025.
- T5.** **Saad-Falcon, J.** *Weaver: Shrinking the Generation-Verification Gap with Weak Verifiers*. Together AI. Jul. 2025.
- T6.** **Saad-Falcon, J.** *Archon: An Architecture Search Framework for Inference-Time Techniques*. LLMs Meet Data Processing Workshop, UC Berkeley. May 2025.
- T7.** **Saad-Falcon, J.** *Archon: An Architecture Search Framework for Inference-Time Techniques*. Databricks. Feb. 2025.
- T8.** **Saad-Falcon, J.** *Archon: An Architecture Search Framework for Inference-Time Techniques*. Contextual AI. Oct. 2024.

Teaching

- Autumn 2025 **Teaching Assistant**, CS 224V: Conversational Virtual Assistants with Deep Learning, Stanford University
- Winter 2025 **Teaching Assistant**, CS 329A: Self-Improving Agents, Stanford University

Mentorship

Current Mentees

2024–present Hangoo Kang

Hannah Gao
Harsh Singh
Matthew Hart
Orhun Akengin
Tanvir Bhathal
Tarun Suresh

Past Mentees

Adrian Lafuente Gamarra; now at Salesforce
Brendan McLaughlin; now at Reflection AI
Herumb Shandilya; now at Mixed Bread
Robby Manihani; now at Pace
Wes Griffin; now at Stanford

Academic Service

2025	Volunteer, Conference on Neural Information Processing Systems (NeurIPS)
2024	Volunteer, Conference on Empirical Methods in Natural Language Processing (EMNLP)
2024	Volunteer, International Conference on Machine Learning (ICML)
2022	Volunteer, Conference of the North American Chapter of the ACL (NAACL)
2020	Volunteer, Conference on Empirical Methods in Natural Language Processing (EMNLP)
2020	Volunteer, Annual Meeting of the Association for Computational Linguistics (ACL)

Leadership & Community Engagement

2025–present	Chair for Computer Science , SoE Dean’s Graduate Student Advisory Council (DGSAC), Stanford University
2021–2022	Co-founder and Student Director , Georgia Tech Venture Capital Fund (Atlanta, GA, USA)
2020–2022	Founding Member and Volunteer , Georgia Tech Computer Science Outreach Club (Atlanta, GA, USA)

Other Experience

2020	Research Analyst , Goldman Sachs, Global Investment Research (New York, NY, USA) Covered the energy and utilities sectors, working with quantitative researchers and traders.
------	---

Last updated: April 2026